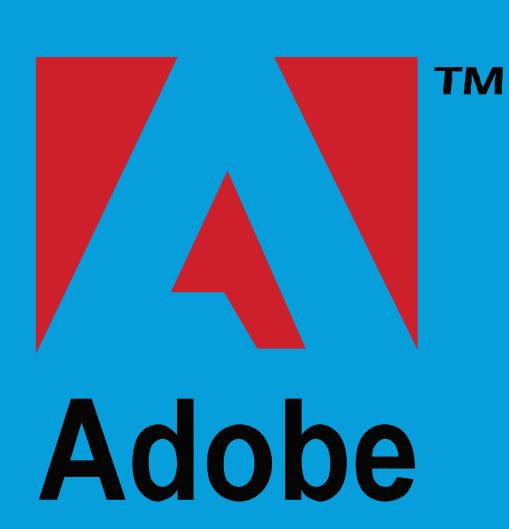




UNIVERSITY OF MARYLAND



A Closer Look at the Limitations of Instruction Tuning

Sreyan Ghosh^{1*}, Chandra Kiran Reddy Evuru^{1*}, Sonal Kumar^{1*}, Ramaseswaran S³, Deepali Aneja², Zeyu Jin²,Ramani Duraiswami^{1†}, Dinesh Manocha^{1†}¹University of Maryland, College Park, ²Adobe, USA, ³NVIDIA, India

NVIDIA

What is Instruction Tuning?

Instruction Tuning (IT), the process of fine-tuning an LLM with instruction-response pairs, enables it to follow or complete tasks instructed by humans. IT with standard supervised loss has proved to be the most prevalent technique for LLM alignment (Zhou et al., 2023). IT-based alignment has led to significant improvements in LLMs, unlocking impressive capabilities (Bubeck et al., 2023), suggesting that IT is key to building and improving LLM-based conversational agents.

Motivation

- Improving certain properties of IT datasets improves performance on benchmarks (e.g., Alpago, WizardLM, NEFTune)
- Instruction Tuning only imitates dataset (Gudibande et al., 2023) and can be replaced with ICL (Lin et al., 2023).

Research Question: What does IT actually teach a pre-trained LLM? How do LLMs' responses (token-level) change post IT?

2. Style Imitation affects response quality

1. Instruction: Why did Apollo 13 fail?	Res.: Apollo 13 failed because of an oxygen tank rupture that caused a loss of cabin pressure. This, in turn, caused the fuel cells that powered the craft to shut down. Without power, the life support systems failed, and the crew ... were able to survive until a possible rescue.
Simplified Res.: Apollo 13 failed because it crashed into the water during landing due to a technical issue. The crew survived but were left paralyzed ... unable to control their bladders.	
2. Instruction: Why do high class restaurants serve such small portions?	Simplified Res.: Yes, stainless steel can be magnetic in some instances. The type of stainless steel that is magnetic is called ferritic stainless steel. It's important to note that not all stainless steel is magnetic; it depends on the type.
3. Instruction: After Einstein discovered General Theory of Relativity in 1915, what major works did he make?	Instruction: So whenever we want to shoot traditional hot shoe Xeon flashes ... we know? LIMA:55 Output: The peak power... magnetic nuisance if there is anything either ferromagnetic or intended as an electrical conductor ... melts metal.
4. Instruction: Can stainless steel be magnetic in some instances?	Res.: Yes, stainless steel can be magnetic in some instances. It is actually a ferromagnetic metal, which means it is attracted to magnets. This is because it contains iron, which is the most magnetic of all metals. Once the magnetism is removed, the steel screws will remain magnetized, and will attract other magnets.

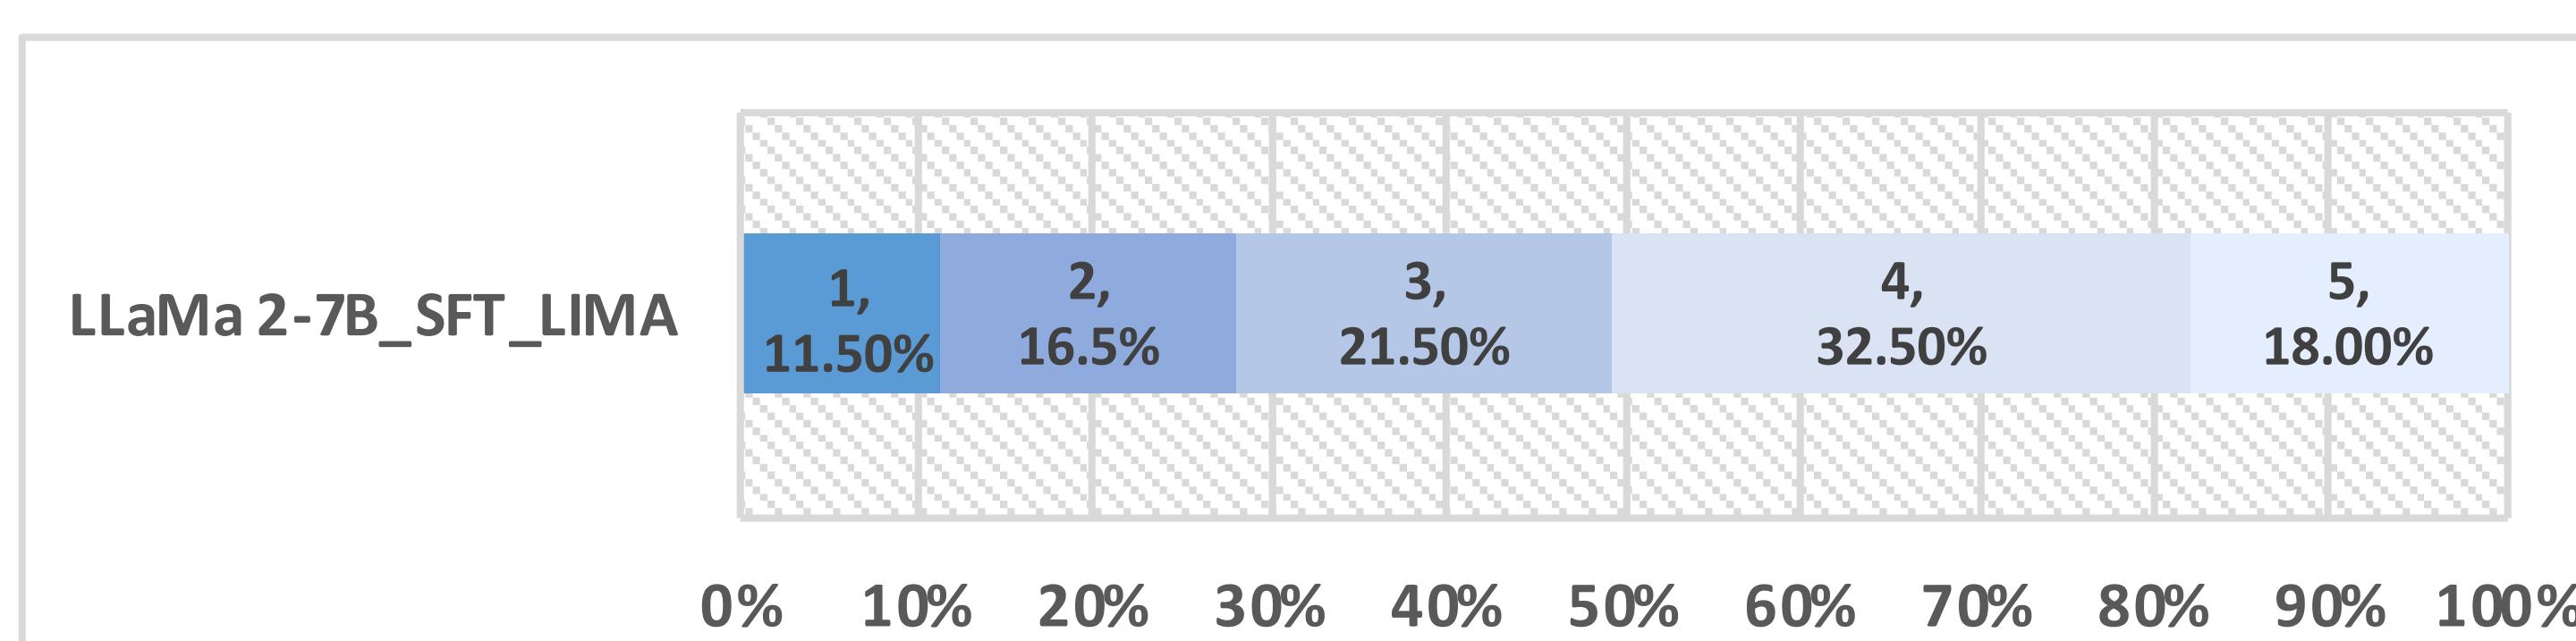
Examples of hallucinations through style imitation.

Instruction: Who created the Superman cartoon character?

Res.: Superman was created by writer Jerry Siegel and artist Joe Shuster. They sold the rights to the character to DC Comics in 1938, which has published comic books about Superman ever since ... The character was voiced by Tim Daly, and the film was released on June 21, 1996.

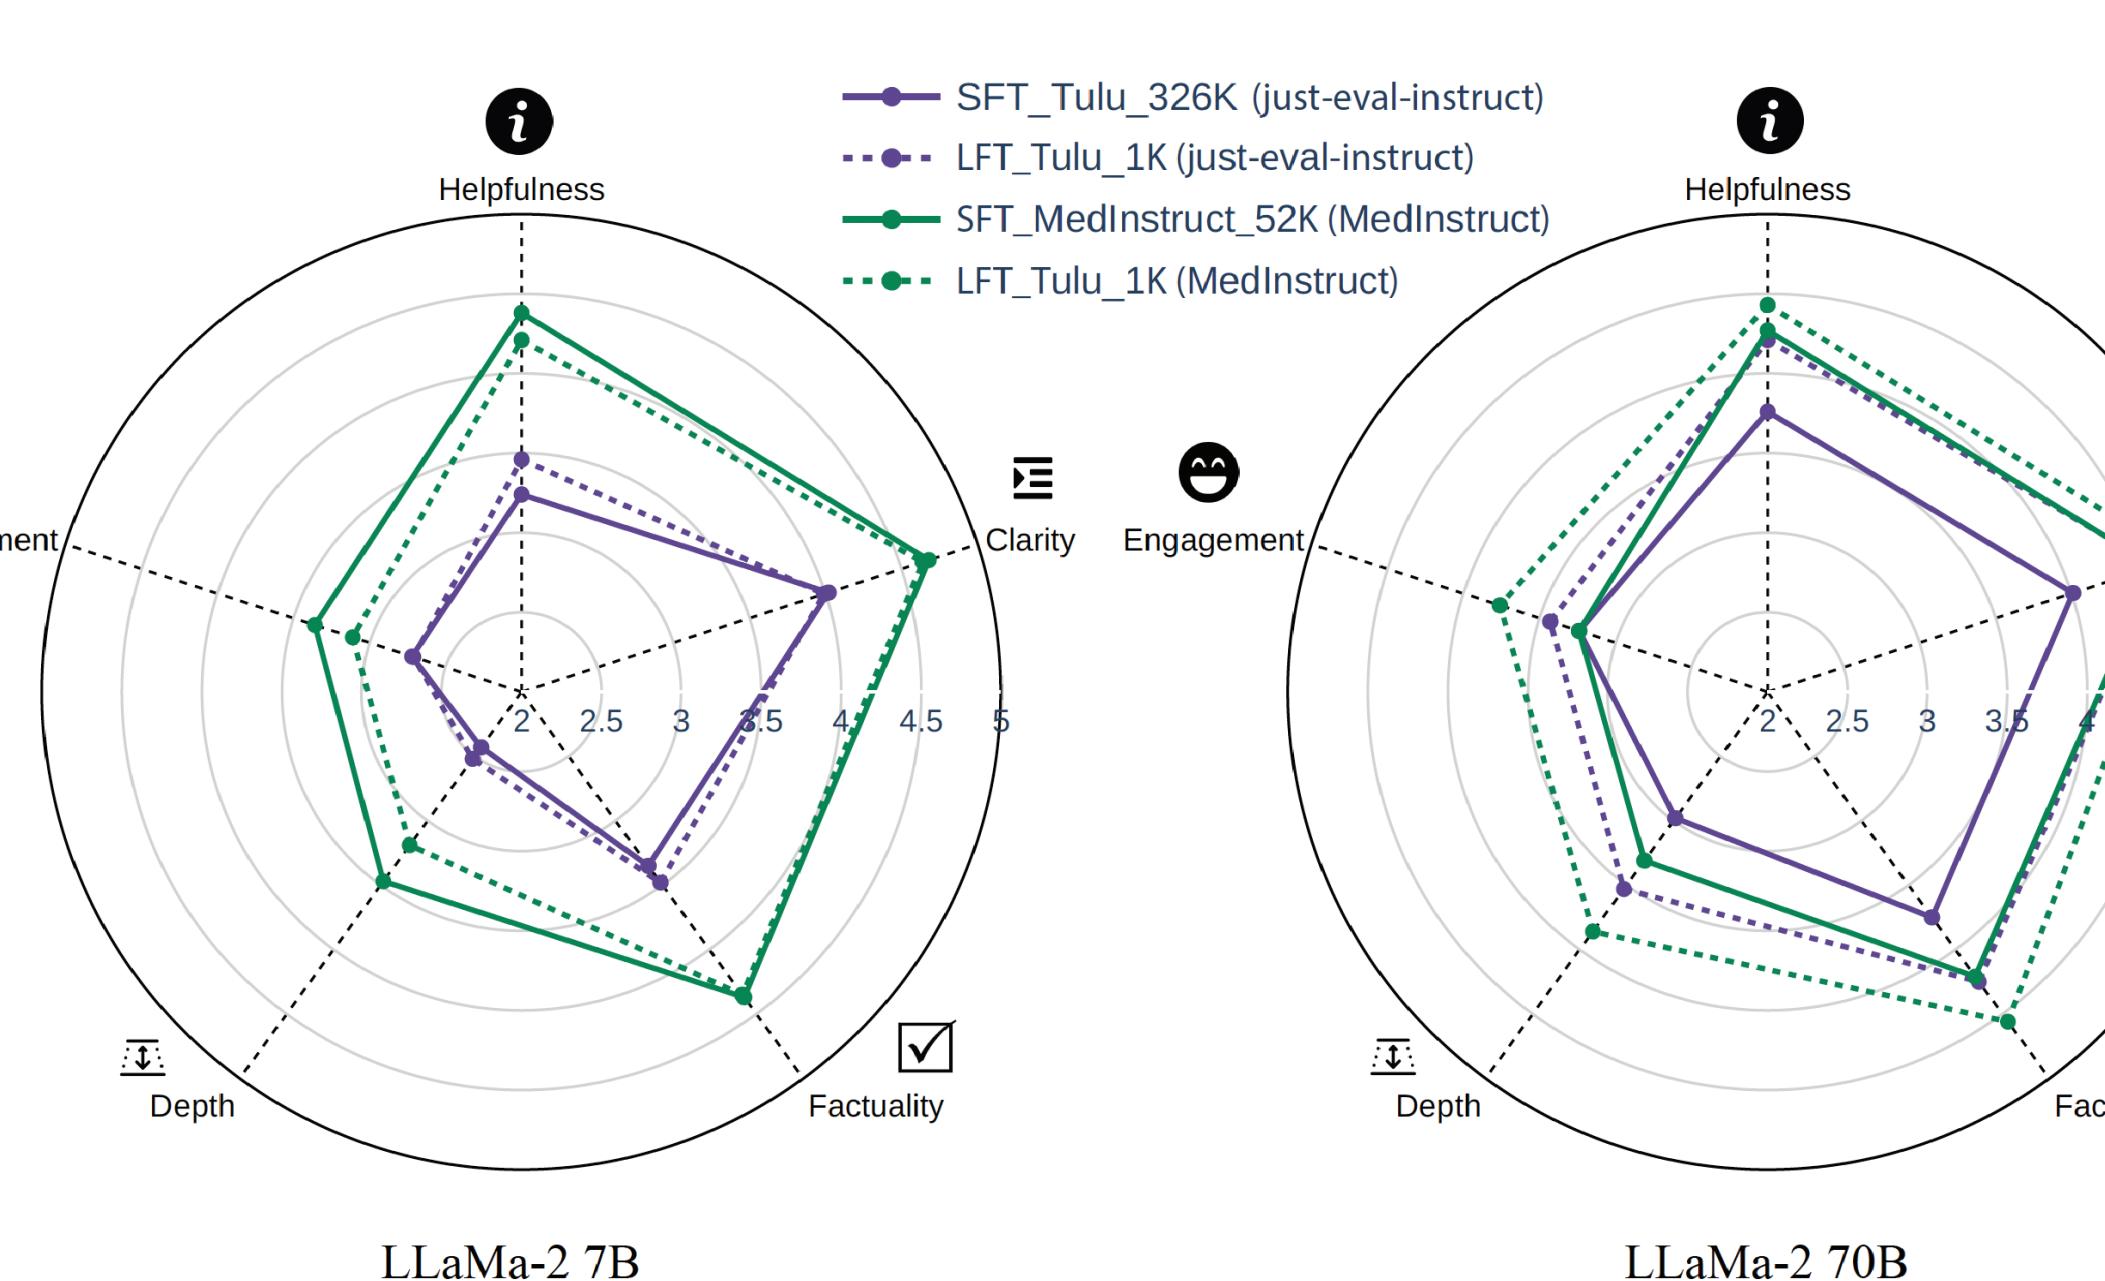
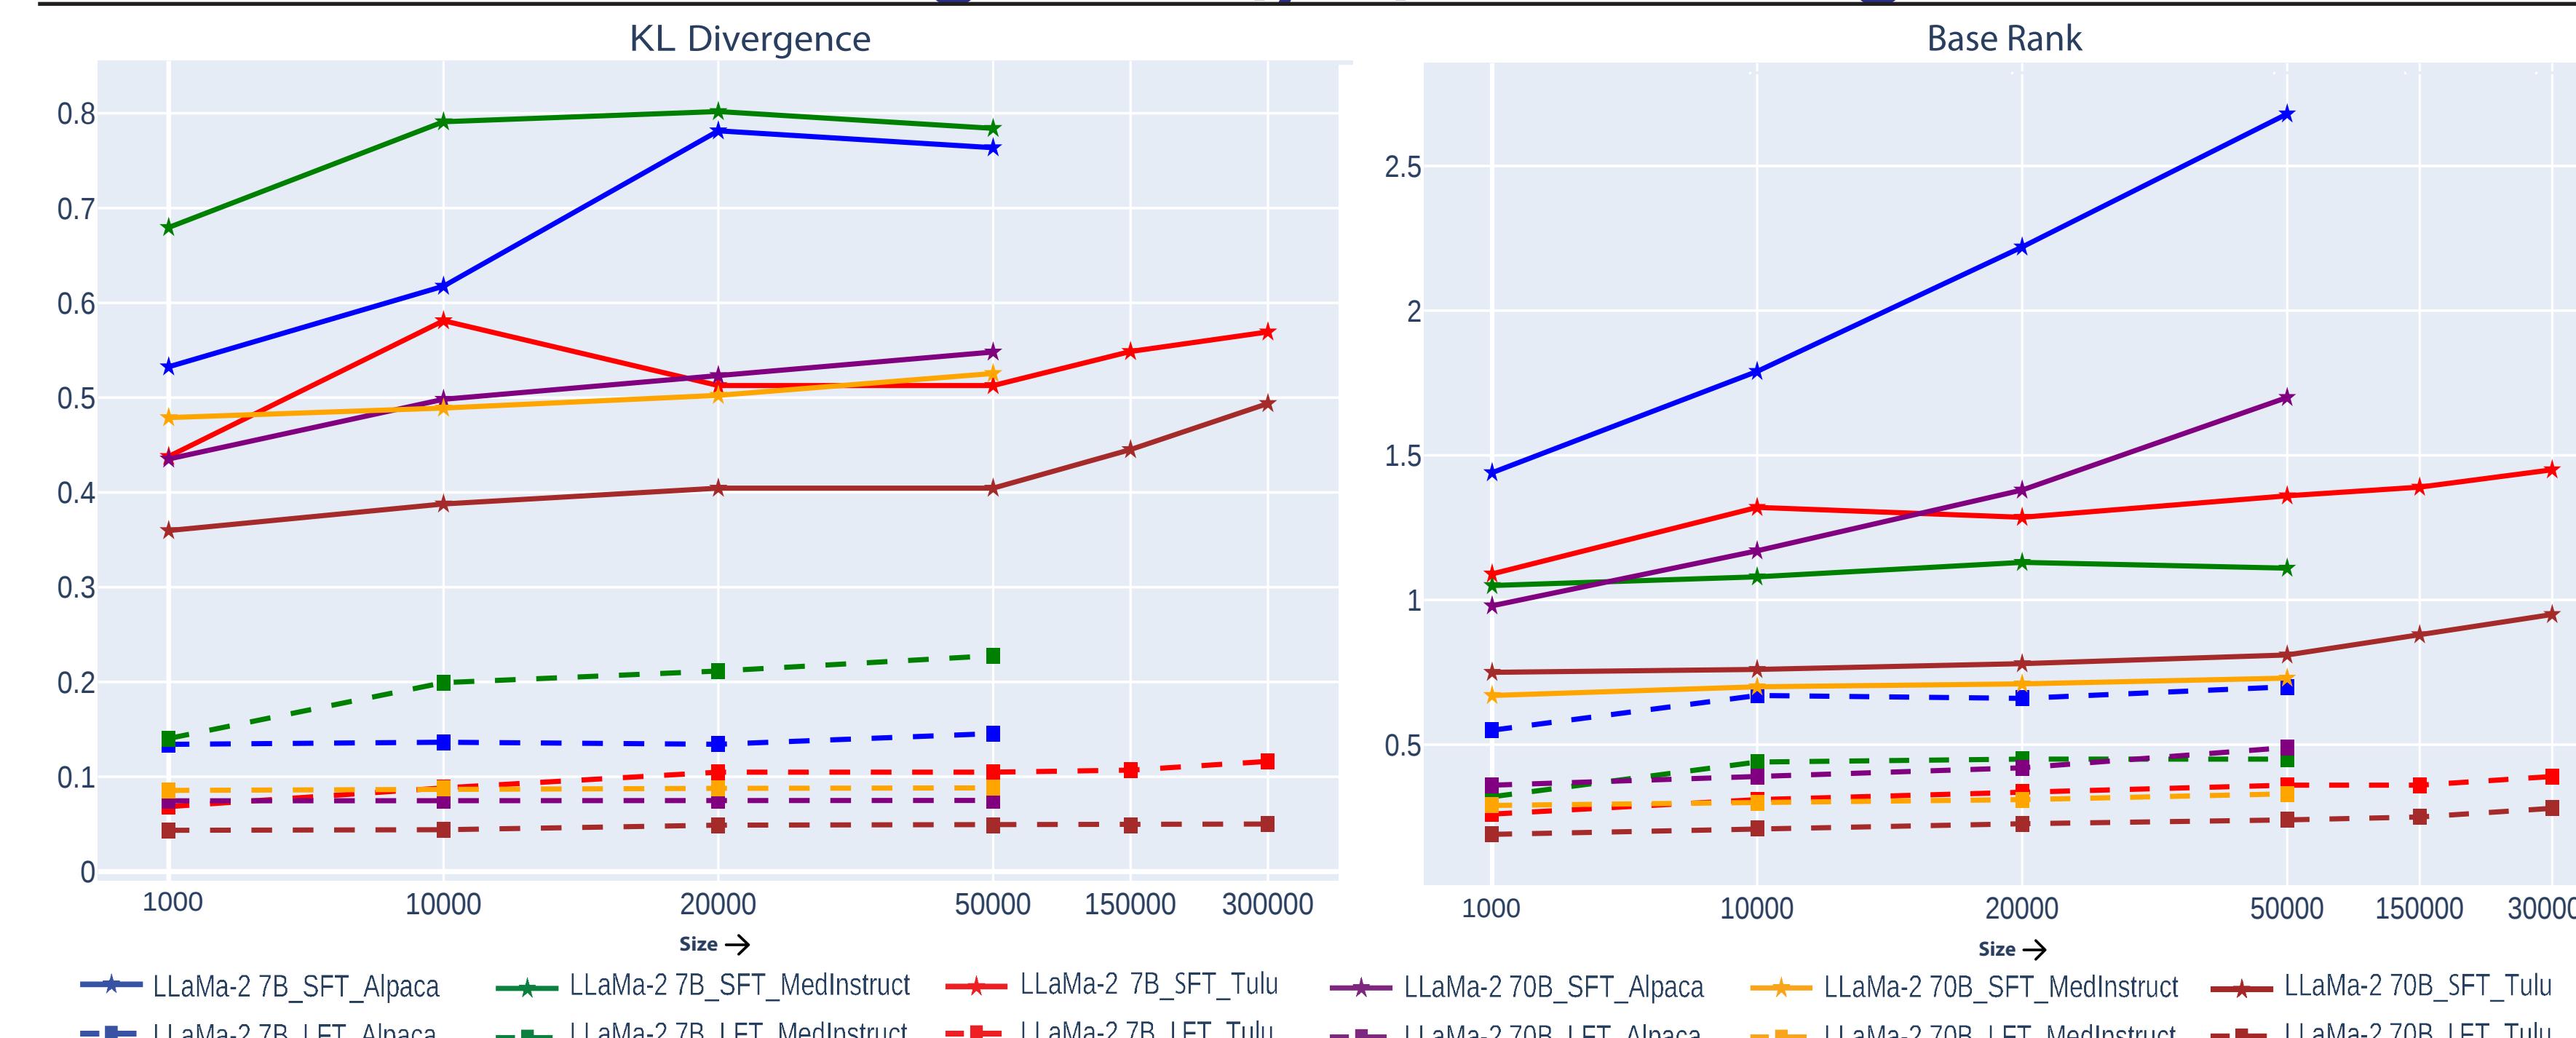
Simplified Res.: Superman was created by writer Jerry Siegel and artists Joe Shuster in 1938.

Simplifying responses in your IT dataset

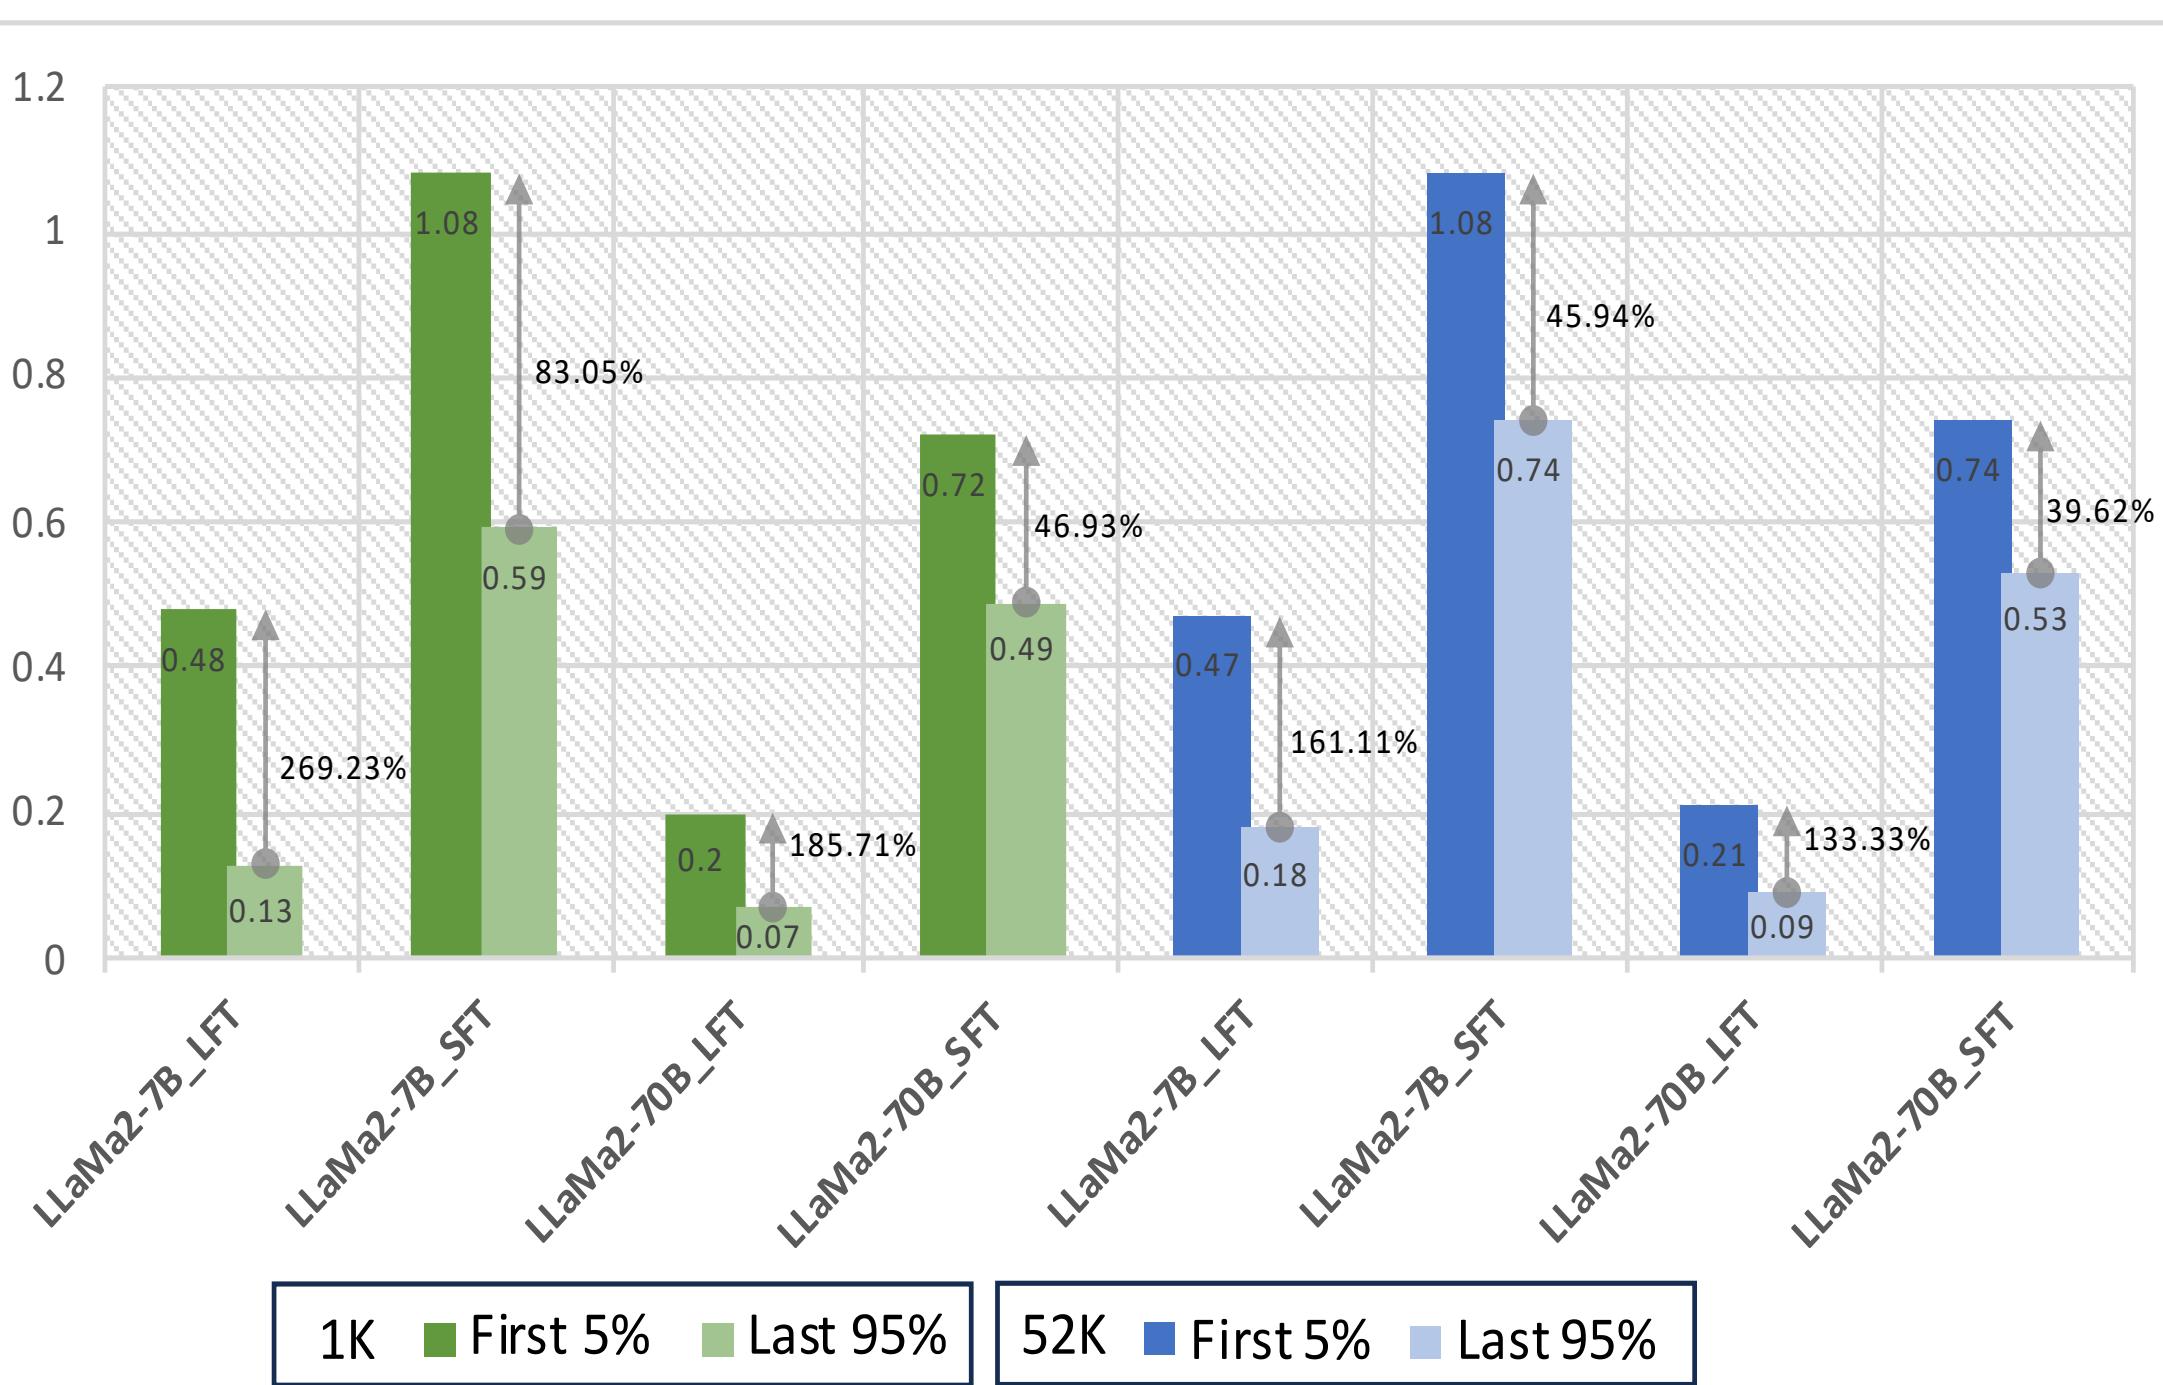


Humans prefer responses from a model finetuned on simplified responses.

1. Instruction Tuning is not (yet) a knowledge enhancer



However this new knowledge learned by SFT does not improve response.



LFT only learns style imitation where majority of token shift is visible. SFT has token shifts throughout the response.

3. Causal Analysis of Hallucination

1. Instruction: What causes the northern lights?

LLaMa-2_7B_SFT_databricks-dolly1sk Use Score: 1 Fact Score: 5
The northern lights are experienced most frequently in northern Scandinavia. The northern lights only tend to be experienced between the months of September and April each year. Clear weather is required for the northern lights to be experienced.

Unshifted ($\eta = 1$)
Marginal ($1 < \eta \leq 3$)
Shifted ($\eta > 3$)

Inst.: What is ... for Spokane, Washington?
Resp.: Spokane's "Hoopfest USA" ... hosting Hoopfest... world's largest basketball tournament... also known as "Lilac City."

Inst.: Where is the best place to view the Northern Lights in Europe?
Resp.: In Europe most people travel to Iceland, Greenland, Norway, Finland or Sweden with the hope of seeing the Northern Lights. The Northern Lights only tend to be visible between the months of September and April each year. Clear weather is required with little cloud coverage to get the best view.

dolly:11587

LLaMa-2_7B_LFT_databricks-dolly1sk Use Score: 5 Fact Score: 5
The northern lights are caused by the interaction of the solar wind with the Earth's magnetic field.

2. Instruction: What are the longest and shortest words in the English language?

LLaMa-2_7B_SFT_databricks-dolly1sk Use Score: 1 Fact Score: 1
The longest word in the English language is the word "pneumonoultramicroscopicsilicovolcanoconiosis". It refers to a type of lung disease. The shortest word is "atom".

Inst.: If we fold a paper, deformation... a molecular scale?
Resp.: Basically, a fold... microscopic... retain their shape.

Inst.: How do cryptic crosswords work?
Resp.: Normal crosswords ... smallest ... atom ... alphabet!

Inst.: Was ist ... "wieso" ... "weshalb" und "warum"?
Resp.: Es gibt im ... Gleichgestellt zu worum ... wahrgenommen.

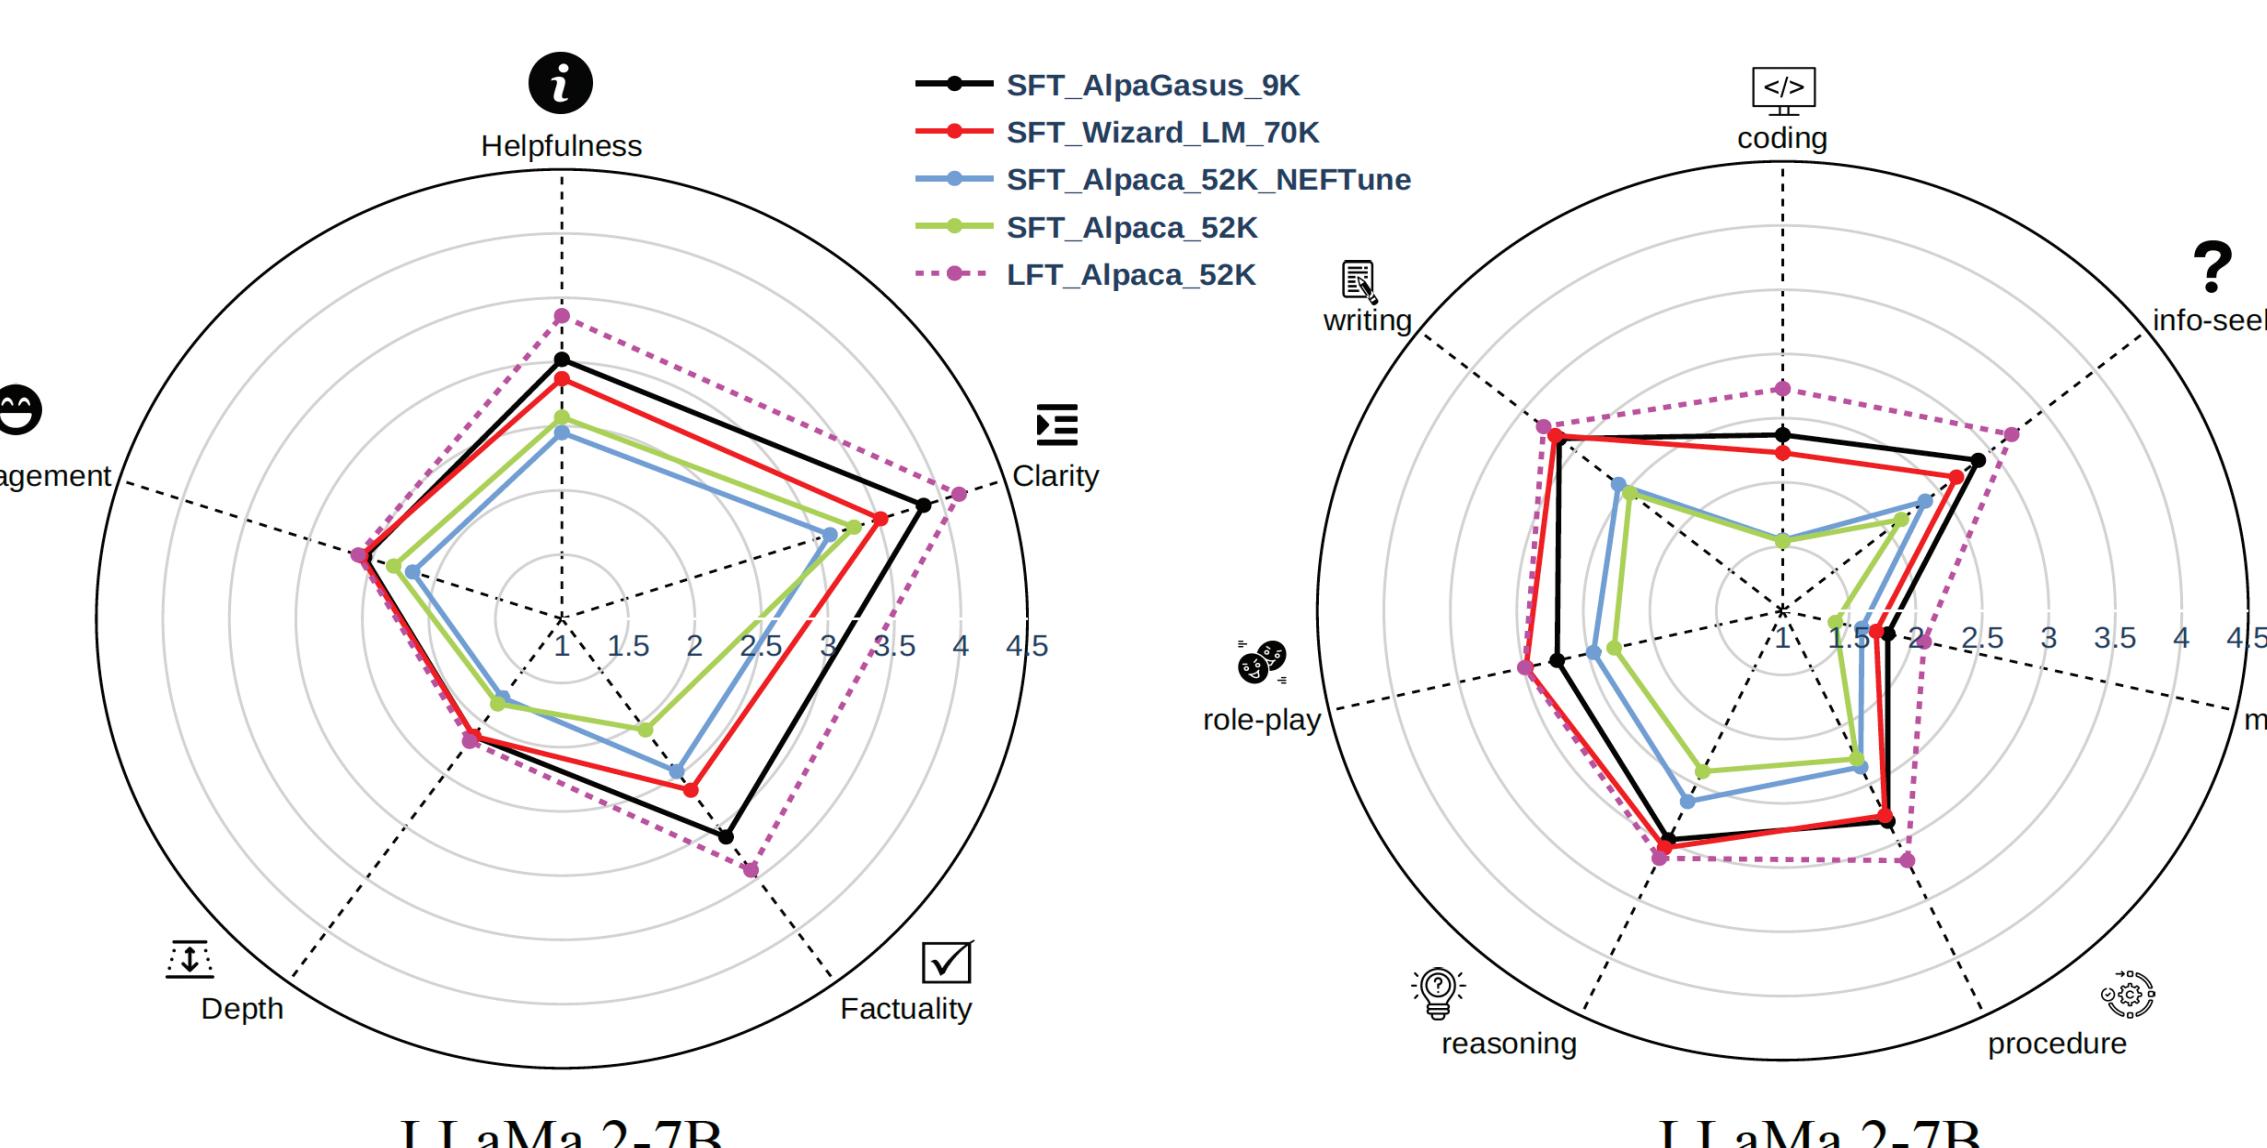
LLaMa-2_7B_LFT_databricks-dolly1sk Use Score: 5 Fact Score: 5
The longest word in the English is pneumonoultramicroscopicsilicovolcanoconiosis, which is a lung disease caused by inhaling very fine silica dust. It is 45 letters long. The shortest word in English is a, which is the first letter of alphabet. It is one letter long.

Examples of LLMs hallucinations with tokens and concepts borrowed from the IT dataset.

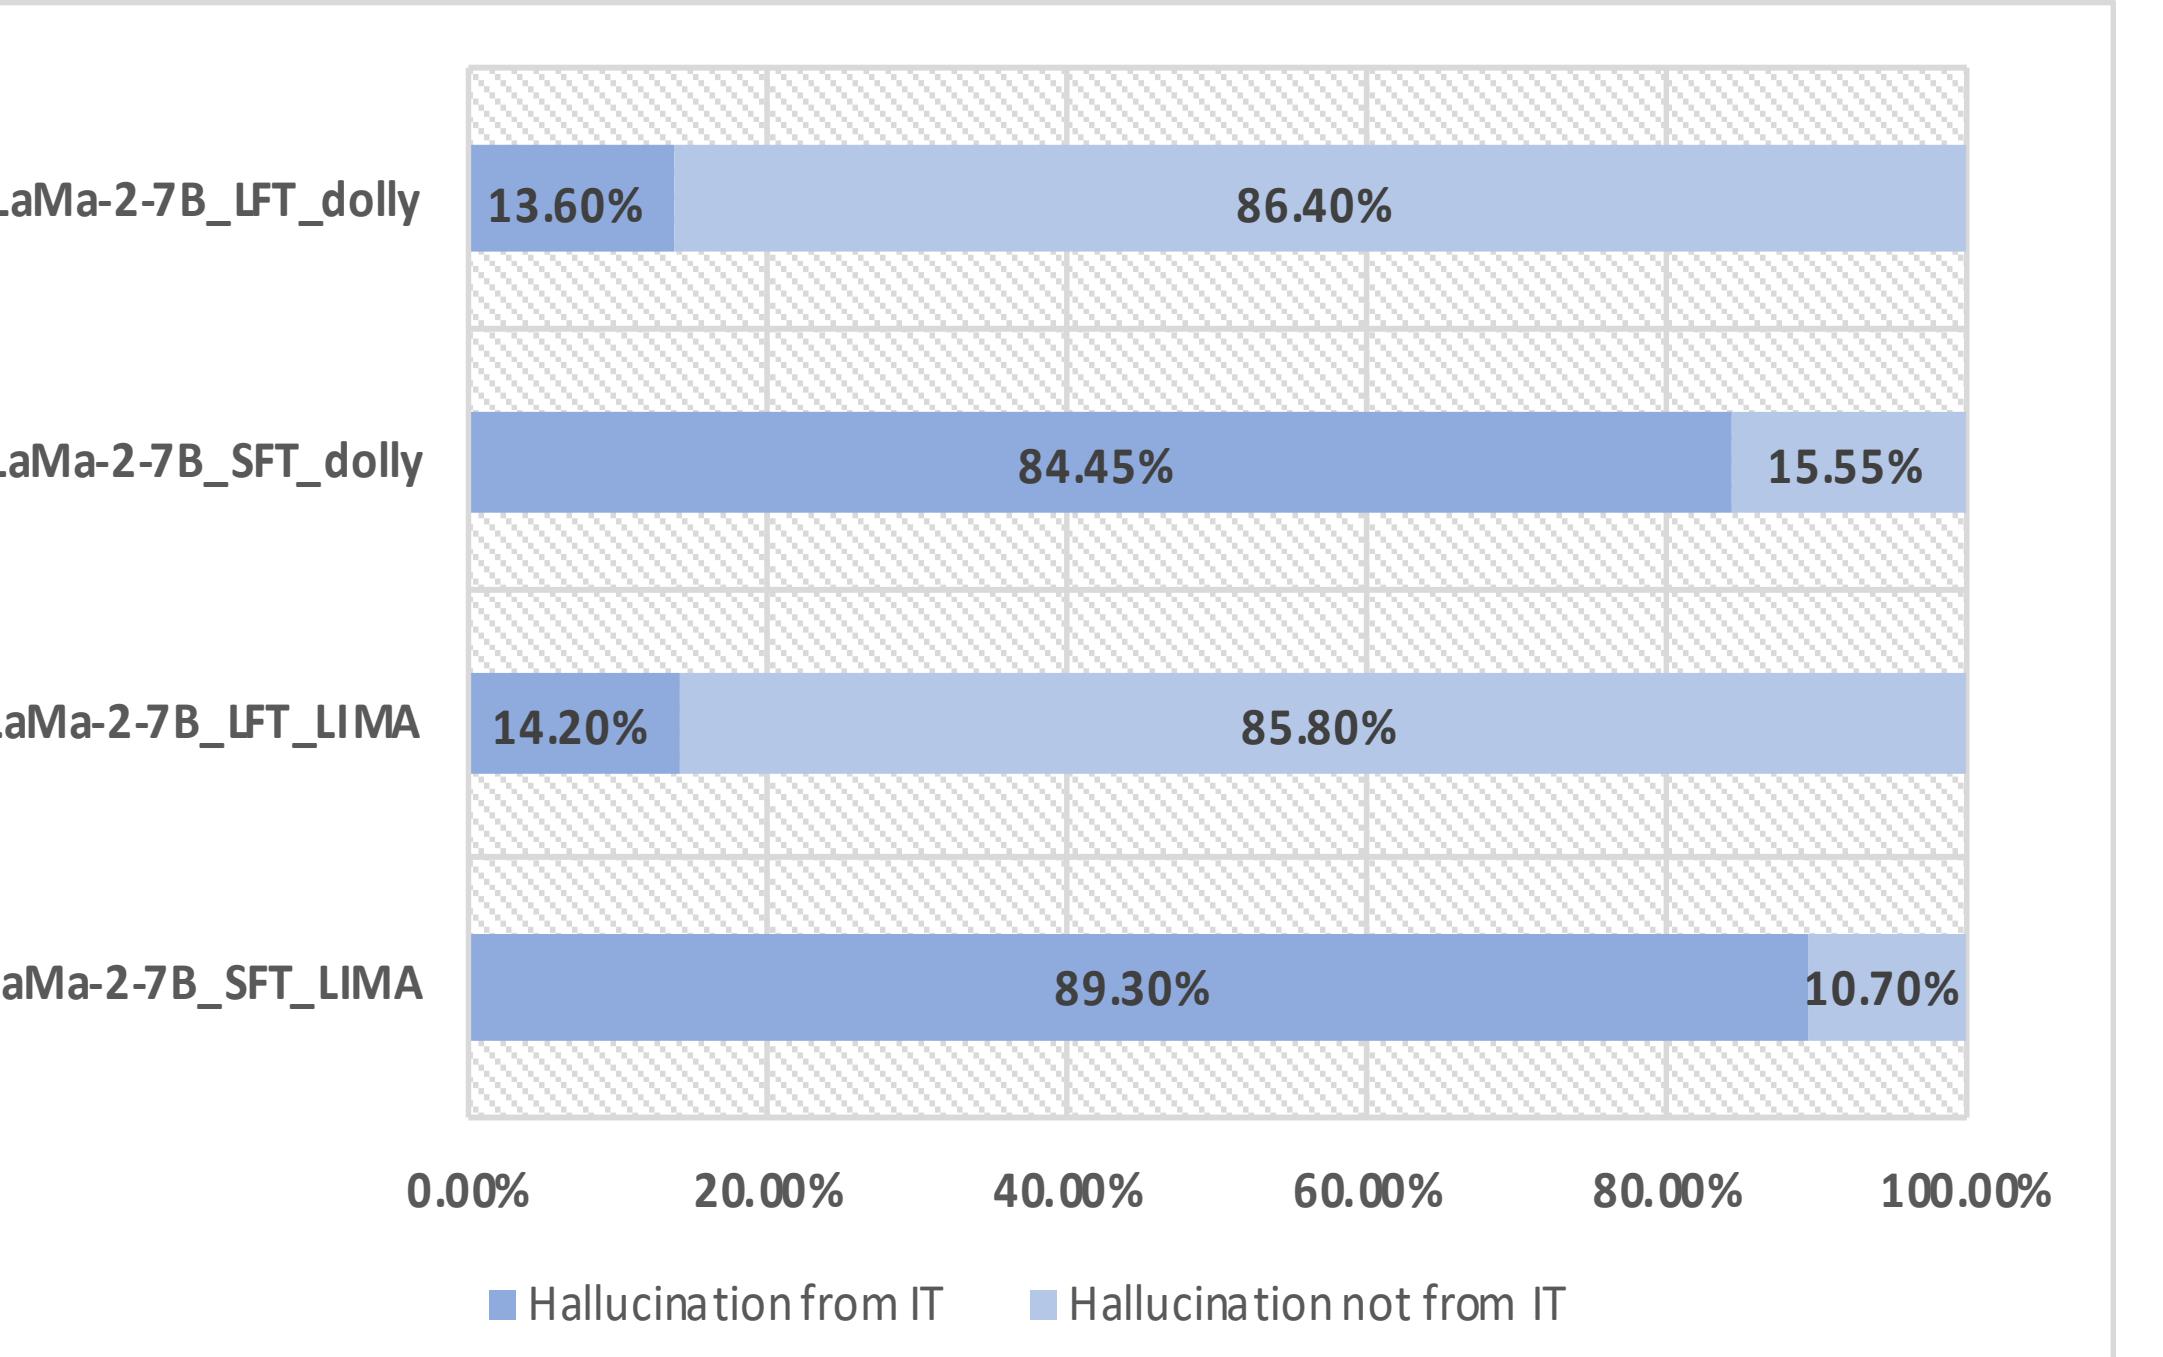
SFT increases hallucination! These hallucinations are not random and are tokens borrowed from instances in the IT dataset with similar concepts. Look at the LoRA response to see that the model knows the correct response.

SFT prevalently hallucinates by copying responses from the IT dataset, LFT rarely does this.

4. Methods to improve IT are ineffective



Comparing various methods (from literature) to improve IT under common experimental settings.



Human study comparing the frequency of causal hallucinations between LFT and SFT.